

MANDY HATHAWAY

AI Ethics & Governance Consultant · Technical Writer · Journalist

mandyhathaway@proton.me · (651) 373-2557 · Remote

Portfolio: mandyhathaway.com · **GitHub:** github.com/mandyhathaway · **Hugging Face:** huggingface.co/philosophyFire

PROFESSIONAL SUMMARY

Independent AI ethics and governance consultant. MA in Ethical Technology and Artificial Intelligence, 30+ years of technology experience, including a decade focused on machine learning and AI, from classical ML through LLM development and fine-tuning. Current work includes governance documentation, vendor evaluation, risk framing, and long-form technical writing for AI builds. Author of the Practical AI Guide for Nonprofits, a vendor-scrutiny framework for mission-driven organizations. Published on algorithmic bias, anthropomorphism, and responsible AI. Available for senior contract roles in AI ethics consulting, governance advisory, vendor evaluation, and technical writing.

CORE SKILLS

AI & ML	Machine Learning, Deep Learning, Neural Networks, Large Language Models, Generative AI, NLP, Computer Vision
Governance	AI Ethics, AI Safety, Responsible AI, Algorithmic Accountability, AI Governance, Vendor Evaluation, Risk Framing
Specialisms	Prompt Engineering, RLHF, LLM Fine-tuning, Bias Evaluation, Model Drift Monitoring, Compliance Documentation
Technical	Python, PyTorch, Hugging Face, NumPy, Pandas, Matplotlib, JavaScript, HTML/CSS, Git, Linux, Azure, Power BI
Writing	Technical Writing, Curriculum Design, Instructional Design, Documentation, Science Writing, Journalism
Leadership	Team Leadership (35+ reports), Hiring & Training, Curriculum Development, Project Management, Public Speaking

EXPERIENCE

AI Ethics & Governance Consultant (Independent)

Mandy Hathaway Consulting · Remote · 2025 to Present

- Phase 2 Documentation Guardianship on Zeki Data AI governance platform: drafting compliance documentation for EU GDPR Article 22 automated decision-making and UK Equality Act obligations on an investigations product.
- Authored Practical AI Guide for Nonprofits, a long-form vendor-scrutiny framework for mission-driven organizations evaluating AI tools, published as open reference.
- Published 95-pair RLHF preference annotation dataset on Hugging Face across six ethics categories (refusal edge cases, sycophancy, parasocial attachment, anthropomorphism, bias, dual-use harm) with accompanying DPO fine-tuning notebook (phi-2, LoRA, Google Colab).

Remote Coding Instructor & Team Lead

Coding With Kids · Remote · 2021 to Present

- Lead remote coding classes for students aged 4 to 17 across Scratch, Python, web development, JavaScript, and Roblox modding via Zoom and Google Meet.
- Served as Team Lead, mentoring and supervising up to 35 instructors; contributed to curriculum development and instructor training standards.
- Qualified for every available performance, retention, and rating-based bonus; recognized as Teacher of the Month (2023).

Technology Editor

The Metropolitan · Saint Paul, MN · 2018 to 2021

- Directed all technology coverage for print and online publication; managed editorial standards and approved all image content.
- Regular contributing writer covering technology, the arts, and social justice; developed ethical standards training for writers and photographers.
- Led hiring, onboarding, and mentoring of new contributors.

Personal Care Assistant

MRCI Worksource · Mankato, MN · 2016 to 2021 · Concurrent with graduate studies

- Full daily care for an adult male quadriplegic client.

Technology Instructor

Science Museum of Minnesota · Saint Paul, MN · 2011 to 2013

- Designed and delivered curriculum in stop motion animation, STEM topics, video editing, and creative technologies for youth and adult learners.
- Taught across onsite and community locations; adapted instruction for diverse audiences and age groups.

EDUCATION

Master of Arts, Ethical Technology and Artificial Intelligence · 3.97 GPA

Metropolitan State University, Saint Paul, MN · 2018 to 2021

Bachelor of Arts, Philosophy (Ethics and Language)

Metropolitan State University, Saint Paul, MN · 2014 to 2017

CERTIFICATIONS

- 2026** AI for Nonprofits Professional Certificate · Microsoft Elevate, LinkedIn & NetHope
- 2026** Operationalizing AI Governance: Strategy and Foundations · All Tech Is Human
- 2026** Hands-On Data Annotation: Applied Machine Learning · LinkedIn
- 2026** Data Preparation, Feature Engineering, and Augmentation for AI Models · LinkedIn
- 2024** Generative AI for Educators · Google
- 2024** Lightweight AI with Small Language Models · LinkedIn

SELECTED PROJECTS

AI Ethics Preference Annotation Dataset

Published RLHF preference annotation dataset on Hugging Face

(huggingface.co/philosophyFire/Ai_ethics_dataset). 95 prompts, 190 response pairs across 6 ethics categories: refusal edge cases, sycophancy, parasocial attachment, anthropomorphism, bias, and dual-use harm. Five scoring dimensions per pair, named failure mode, and written justification for every dispreferred response. Includes a DPO fine-tuning notebook built in Google Colab using microsoft/phi-2 with LoRA adapters. Published CC BY 4.0.

PUBLICATIONS

The Bias We're Building In · Essay on algorithmic bias and responsible AI development

We Built the Monster We're Afraid Of · Essay on AI anthropomorphism and projected fears, 2020

Tools for Trusting Information in a World of Information Overload · The Metropolitan, 2021

Data versus Lore · AI journalism series, The Metropolitan, 2020